TRƯỜNG ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

Hadoop distributed file system (HDFS)

Instructor: Dr. Thanh-Chung Dao

Slides by Dr. Viet-Trung Tran

School of Information and Communication Technology

1

Overview of HDFS

- Provides inexpensive and reliable storage for massive amounts of data
- Designed for
  - Big files (100 MB to several TBs file sizes)
  - Write once, read many times (Appending only)
  - Running on commodity hardware
- Hierarchical UNIX style file systems
  - (e.g., /hust/soict/hello.txt)
 - UNIX style file ownership and permissions

2

2

HDFS main design principles

- I/O pattern
  - Append only → reduce synchronization
- Data distribution
  - File is splitted in big chunks (64 MB)
  → reduce metadata size
  → reduce network communication
- Data replication
  - Each chunk is usually replicated in 3 different nodes
- Fault tolerance
  - Data node: re-replication
  - Name node
    - Secondary Namenode
    - Standby, Active Namenodes

3

HDFS Architecture

- Master/slave architecture
- HDFS master: Namenode
  - Manage namespace and metadata
  - Monitor Datanode
- HDFS slaves: Datanodes
  - Handle read/write the actual data {chunks}
  - Chunks are local files in the local file systems

4

Functions of a Namenode

- Manages File System Namespace
  - Maps a file name to a set of blocks
  - Maps a block to the Datanodes where it resides
- Cluster Configuration Management
- Replication Engine for Blocks

5

Namenode metadata

- Metadata in memory
  - The entire metadata is in main memory
  - No demand paging of metadata
- Types of metadata
  - List of files
  - List of Blocks for each file
  - List of Datanodes for each block
  - File attributes, e.g. creation time, replication factor
- A Transaction Log
  - Records file creations, file deletions etc

6

Datanode

- A Block Server
  - Stores data in the local file system (e.g. ext3)
  - Stores metadata of a block (e.g. CRC)
  - Serves data and metadata to Clients
- Block Report
  - Periodically sends a report of all existing blocks to the Namenode
- Facilitates Pipelining of Data
  - Forwards data to other specified Datanodes
- Heartbeat
  - Datanodes send heartbeat to the Namenode
    - Once every 3 seconds
  - Namenode uses heartbeats to detect Datanode failure

7

Data replication

- Chunk placement
  - Current Strategy
    - One replica on local node
    - Second replica on a remote rack
    - Third replica on same remote rack
    - Additional replicas are randomly placed
  - Clients read from nearest replicas
- Namenode detects Datanode failures
  - Chooses new Datanodes for new replicas
  - Balances disk usage
  - Balances communication traffic to Datanodes

8

Data rebalance

- Goal: % disk full on Datanodes should be similar
  - Usually run when new Datanodes are added
  - Cluster is online when Rebalancer is active
  - Rebalancer is throttled to avoid network congestion
 - Command line tool

9

Data correctness

- Use Checksums to validate data
  - Use CRC32
- File Creation
  - Client computes checksum per 512 bytes
  - Datanode stores the checksum
- File access
  - Client retrieves the data and checksum from Datanode
  - If Validation fails, Client tries other replicas

10

Data pipelining

- Client retrieves a list of Datanodes on which to place replicas of a block
- Client writes block to the first Datanode
- The first Datanode forwards the data to the next node in the Pipeline
- When all replicas are written, the Client moves on to write the next block in file
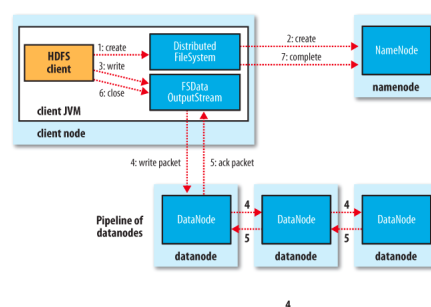

11

Secondary Name node

- Namenode is a single point of failure
- Secondary Namenode
  - Checkpointing latest copy of the FsImage and the Transaction Log files.
    - Copies FsImage and Transaction Log from Namenode to a temporary directory
- When Namenode restarted
  - Merges FSImage and Transaction Log into a new FSImage in temporary directory
  - Uploads new FSImage to the Namenode
    - Transaction Log on Namenode is purged

12

Namenode high availability (HA)
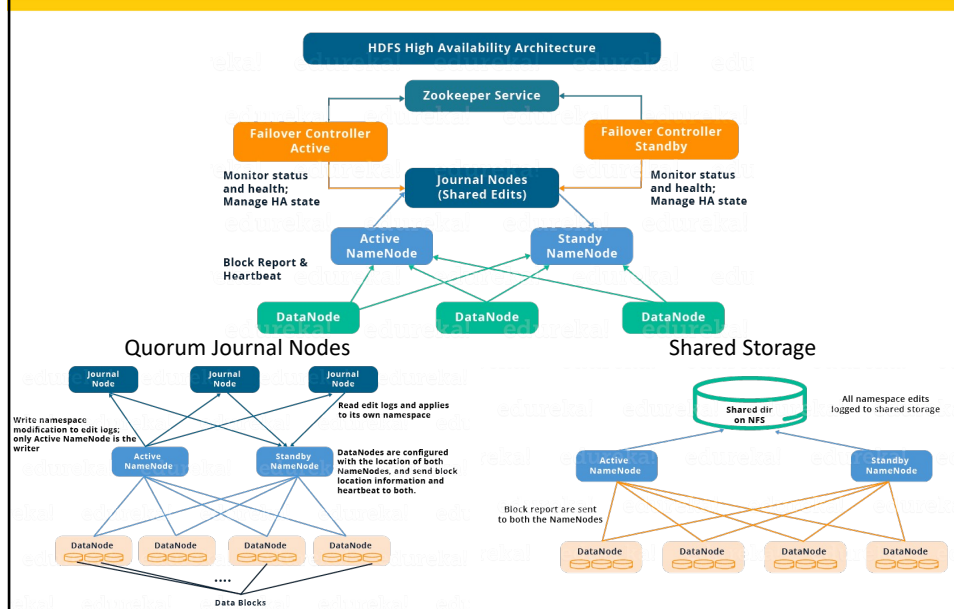
13

HDFS command-line interface

List Files	
hdfs dfs -ls /	List all the files/directories for the given hdfs destination path.
hdfs dfs -ls -d /hadoop	Directories are listed as plain files. In this case, this command will list the details of hadoop folder.
hdfs dfs -ls -h /data	Format file sizes in a human-readable fashion (eg 64.0m instead of 67108864).
hdfs dfs -ls -R /hadoop	Recursively list all files in hadoop directory and all subdirectories in hadoop directory.
hdfs dfs -ls /hadoop/dat*	List all the files matching the pattern. In this case, it will list all the files inside hadoop directory which starts with 'dat'.

Read/Write Files	
hdfs dfs -text /hadoop/derby.log	HDFS Command that takes a source file and outputs the file in text format on the terminal. The allowed formats are zip and TextRecordInputStream.
hdfs dfs -cat /hadoop/test	This command will display the content of the HDFS file test on your stdout .
hdfs dfs -appendToFile /home/ubuntu/test1 /hadoop/text2	Appends the content of a local file test1 to a hdfs file test2.

14

Upload, download files

Upload/Download Files	
hdfs dfs -put /home/ubuntu/sample /hadoop	**Copies the file from local file system to HDFS.**
hdfs dfs -put -f /home/ubuntu/sample /hadoop	**Copies the file from local file system to HDFS, and in case the local already exits in the given destination path, using -f option with put command will overwrite it.**
hdfs dfs -put -l /home/ubuntu/sample /hadoop	**Copies the file from local file system to HDFS. Allow DataNode to lazily persist the file to disk. Forces replication factor of 1.**
hdfs dfs -put -p /home/ubuntu/sample /hadoop	**Copies the file from local file system to HDFS. Passing -p preserves access and modification times, ownership and the mode.**
hdfs dfs -get /newfile /home/ubuntu/	**Copies the file from HDFS to local file system.**
hdfs dfs -get -p /newfile /home/ubuntu/	**Copies the file from HDFS to local file system. Passing -p preserves access and modification times, ownership and the mode.**
hdfs dfs -get /hadoop/*.txt /home/ubuntu/	**Copies all the files matching the pattern from local file system to HDFS.**
hdfs dfs -copyFromLocal /home/ubuntu/sample /hadoop	**Works similarly to the put command, except that the source is restricted to a local file reference.**
hdfs dfs -copyToLocal /newfile /home/ubuntu/	**Works similarly to the put command, except that the destination is restricted to a local file reference.**
hdfs dfs -moveFromLocal /home/ubuntu/sample /hadoop	**Works similarly to the put command, except that the source is deleted after it's copied.**

15

File management

File Management	
hdfs dfs -cp /hadoop/file1 /hadoop1	**Copies file from source to destination on HDFS. In this case, copying file1 from hadoop directory to hadoop1 directory.**
hdfs dfs -cp -p /hadoop/file1 /hadoop1	**Copies file from source to destination on HDFS. Passing -p preserves access and modification times, ownership and the mode.**
hdfs dfs -cp -f /hadoop/file1 /hadoop1	**Copies file from source to destination on HDFS. Passing -f overwrites the destination if it already exists.**
hdfs dfs -mv /hadoop/file1 /hadoop1	**Move files that match the specified file pattern <src> to a destination <dst>. When moving multiple files, the destination must be a directory.**
hdfs dfs -rm /hadoop/file1	**Deletes the file (sends it to the trash).**
hdfs dfs -rm -r /hadoop **hdfs dfs -rm -R /hadoop** **hdfs dfs -rmr /hadoop**	**Deletes the directory and any content under it recursively.**
hdfs dfs -rm -skipTrash /hadoop	**The -skipTrash option will bypass trash, if enabled, and delete the specified file(s) immediately.**
hdfs dfs -rm -f /hadoop	**If the file does not exist, do not display a diagnostic message or modify the exit status to reflect an error.**
hdfs dfs -rmdir /hadoop1	**Delete a directory.**
hdfs dfs -mkdir /hadoop2	**Create a directory in specified HDFS location.**
hdfs dfs -mkdir -f /hadoop2	**Create a directory in specified HDFS location. This command does not fail even if the directory already exists.**
hdfs dfs -touchz /hadoop3	**Creates a file of zero length at <path> with current time as the timestamp of that <path>.**

16

Ownership and validation

Ownership and Validation	
hdfs dfs -checksum /hadoop/file1	Dump checksum information for files that match the file pattern <src> to stdout.
hdfs dfs -chmod 755 /hadoop/file1	Changes permissions of the file.
hdfs dfs -chmod -R 755 /hadoop	Changes permissions of the files recursively.
hdfs dfs -chown ubuntu:ubuntu /hadoop	Changes owner of the file. 1st ubuntu in the command is owner and 2nd one is group.
hdfs dfs -chown -R ubuntu:ubuntu /hadoop	Changes owner of the files recursively.
hdfs dfs -chgrp ubuntu /hadoop	Changes group association of the file.
hdfs dfs -chgrp -R ubuntu /hadoop	Changes group association of the files recursively.
Filesystem	
hdfs dfs -df /hadoop	Shows the capacity, free and used space of the filesystem.
hdfs dfs -df -h /hadoop	Shows the capacity, free and used space of the filesystem. -h parameter Formats the sizes of files in a human-readable fashion.
hdfs dfs -du /hadoop/file	Show the amount of space, in bytes, used by the files that match the specified file pattern.
hdfs dfs -du -s /hadoop/file	Rather than showing the size of each individual file that matches the pattern, shows the total (summary) size.
hdfs dfs -du -h /hadoop/file	Show the amount of space, in bytes, used by the files that match the specified file pattern. Formats the sizes of files in a human-readable fashion.

17

Administration

Administration	
hdfs balancer -threshold 30	Runs a cluster balancing utility. Percentage of disk capacity. This overwrites the default threshold.
hadoop version	To check the vesrion of Hadoop.
hdfs fsck /	It checks the health of the Hadoop file system.
hdfs dfsadmin -safemode leave	The command to turn off the safemode of NameNode.
hdfs dfsadmin -refreshNodes	Re-read the hosts and exclude files to update the set of Datanodes that are allowed to connect to the Namenode and those that should be decommissioned or recommissioned.
hdfs namenode -format	Formats the NameNode.

18

HDFS Name node UI

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Overview 'hd01:8020' (active)

Started:	Thu Mar 14 11:01:37 +0700 2019
Version:	3.1.1.3.1.0.0-78, re4f82af51faec922b4804d0232a637422ec29e64
Compiled:	Thu Dec 06 20:34:00 +0700 2018 by jenkins from (HEAD detached at e4f82af)
Cluster ID:	CID-a1dea38d-b6cf-44e4-b5e8-3bde87ffad35
Block Pool ID:	BP-1412866890-10.10.137.41-1544787807355

Summary

Security is off.

Safemode is off.

165,739 files and directories, 99,858 blocks (99,858 replicated blocks, 0 erasure coded block groups) = 265,597 total filesystem object(s).

Heap Memory used 250.76 MB of 1011.25 MB Heap Memory. Max Heap Memory is 1011.25 MB.

19

HDFS Name node UI (2)

In operation

Show 25 ▾ entries   Search:

Node	Http Address	Last contact	Last Block Report	Capacity	Blocks	Block pool used	Version
✔ hd01:50010 (10.10.137.41:50010)	http://hd01:50075	1s	251m	296.83 GB	99858	166.56 GB (56.11%)	3.1.1.3.1.0.0-78
✔ hd02:50010 (10.10.137.42:50010)	http://hd02:50075	0s	57m	296.83 GB	99858	166.57 GB (56.12%)	3.1.1.3.1.0.0-78
✔ hd03:50010 (10.10.137.43:50010)	http://hd03:50075	2s	197m	296.83 GB	99858	166.57 GB (56.12%)	3.1.1.3.1.0.0-78

Showing 1 to 3 of 3 entries   Previous 1 Next

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Browse Directory

| / | | Go! |

Show 25 ▾ entries   Search:

	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
☐	drwxrwxrwt	yarn	hadoop	0 B	Jun 01 17:52	0	0 B	app-logs	🗑
☐	drwxr-xr-x	hdfs	hdfs	0 B	Dec 18 2018	0	0 B	apps	🗑
☐	drwxr-xr-x	yarn	hadoop	0 B	Dec 14 2018	0	0 B	ats	🗑
☐	drwxr-xr-x	hdfs	hdfs	0 B	Dec 14 2018	0	0 B	atsv2	🗑

20

Other HDFS interfaces

- Java API
- Thrift API
- Fuse
- WebDAV

21

HDFS data format

Text
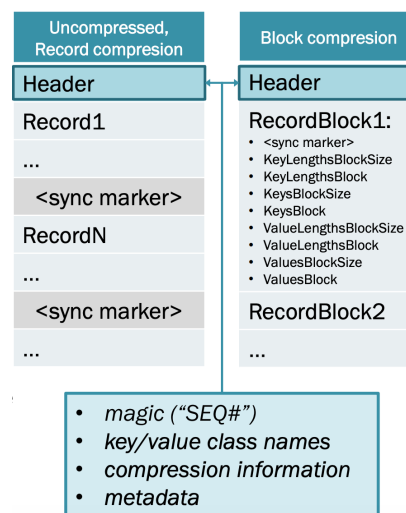Sequence file
Avro
Parquet
Optimized Row Columnar (ORC)

22

Text file

- CSV, TSV, Json records
- Convenient format to use to exchange between applications or scripts
- Human readable and parsable
- Do not support block compression
- Not as efficient to query
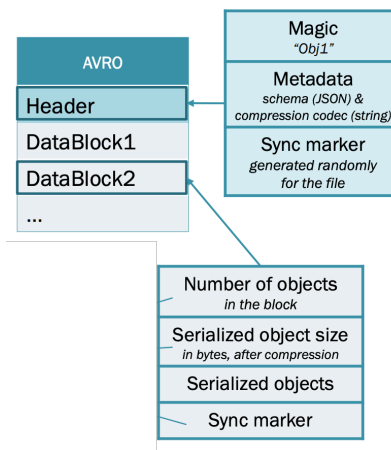- Good for the beginning, but not good enough for real life.

23

Sequence file

- Provides a persistent data structure for binary key-value pairs
- Commonly used to transfer data between Map Reduce jobs
- Can be used as an archive to pack small files in Hadoop
- Row-based
- Compression
- Splittable
  - Support splitting even when the data is compressed

Uncompressed, Record compresion	Block compresion
Header	Header
Record1	RecordBlock1: • <sync marker> • KeyLengthsBlockSize • KeyLengthsBlock • KeysBlockSize • KeysBlock • ValueLengthsBlockSize • ValueLengthsBlock • ValuesBlockSize • ValuesBlock
...	
<sync marker>	
RecordN	
...	
<sync marker>	RecordBlock2
...	...

- *magic ("SEQ#")*
- *key/value class names*
- *compression information*
- *metadata*

24

Avro

- Row based
- Supports (object) compression and splitting
- Flexible data scheme
  - Schema (JSON) included to the file
- Data types
  - primitive: null, boolean, int, long, ...
  - complex: records, arrays, maps, ...
- Binary and JSON data serialization
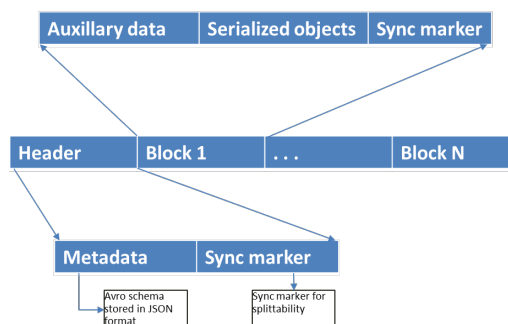- Data corruption detection

AVRO
Header
DataBlock1
DataBlock2
...

Magic
"Obj1"
Metadata
schema (JSON) & compression codec (string)
Sync marker
generated randomly for the file

Number of objects
in the block
Serialized object size
in bytes, after compression
Serialized objects
Sync marker

25

Avro – File structure and example

Sample AVRO schema in JSON format

```
{
  "type" : "record",
  "name" : "tweets",
  "fields" : [ {
    "name" : "username",
    "type" : "string",
}, {
    "name" : "tweet",
    "type" : "string",
}, {
    "name" : "timestamp",
    "type" : "long",
} ],
  "doc:" : "schema for storing tweets"
}
```

Avro file structure

Auxillary data	Serialized objects	Sync marker

Header	Block 1	. . .	Block N

Metadata	Sync marker

Avro schema stored in JSON format

Sync marker for splittability

26

Parquet

- Column-oriented binary file format
- Efficient in terms of disk I/O when specific columns need to be queried
- Supports (page) compression and splitting
- Supports nested columns (Dremel encoding)

Internal structure of parquet file

Nested schema

Table representation

X	Y	Z
x1	y1	z1
x2	y2	z2
x3	y3	z3
x4	y4	z4
x5	Y5	z5

Row format

| x1 | y1 | z1 | x2 | y2 | z2 | x3 | y3 | z3 | x4 | y4 | z4 | x5 | y5 | z5 |

Columnar format

| x1 | x2 | x3 | x4 | x5 | y1 | y2 | y3 | y4 | y5 | z1 | z2 | z3 | z4 | z5 |

encoded chunk   encoded chunk   encoded chunk
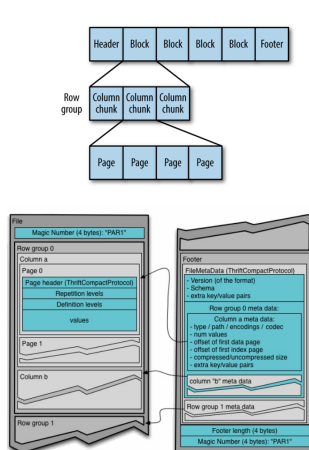
27

Parquet file structure & configuration

Internal structure of parquet file

Configurable parquet parameters

| Property name	Default value	Description
parquet.block.size	128 MB	The size in bytes of a block (row group).
parquet.page.size	1MB	The size in bytes of a page.
parquet.dictionary.page.size	1MB	The maximum allowed size in bytes of a dictionary before falling back to plain encoding for a page.
parquet.enable.dictionary	true	Whether to use dictionary encoding.
parquet.compression	UNCOMPRESSED	The type of compression: UNCOMPRESSED, SNAPPY, GZIP & LZO

Parquet is optimized for high compression and high scan efficiency

28

28

Optimized row columnar (ORC)

- RCFile
  - Every column is compressed individually within the row group
- ORC File
  - Block-mode compression
  - Data type support
  - Ordered data store (within one stripe)
- Stores collections of rows and within the collection the row data is stored in columnar format
- Introduces a lightweight indexing that enables skipping of irrelevant blocks of rows
- Splittable: allows parallel processing of row collections
- Indices with column-level aggregated values (min, max, sum and count)
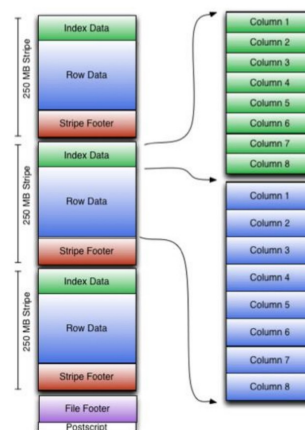
29

TRƯỜNG ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

Thank you for your attention!
Q&A

30