

Parallel Programming with Hadoop/MapReduce

Slides from Tao Yang


1

Overview

- **Related technologies**
 - Hadoop/Google file system
- **MapReduce applications**

2

Typical Hadoop Cluster



- 40 nodes/rack, 1000-4000 nodes in cluster
- 1 Gbps bandwidth in rack, 8 Gbps out of rack
- Node specs :
8-16 cores, 32 GB RAM, 8 × 1.5 TB disks

5

MapReduce Programming Model

- Inspired from map and reduce operations commonly used in functional programming languages like Lisp.
- Have multiple map tasks and reduce tasks
- Users implement interface of two primary methods:
 - Map: $(key1, val1) \rightarrow (key2, val2)$
 - Reduce: $(key2, [val2]) \rightarrow [val3]$

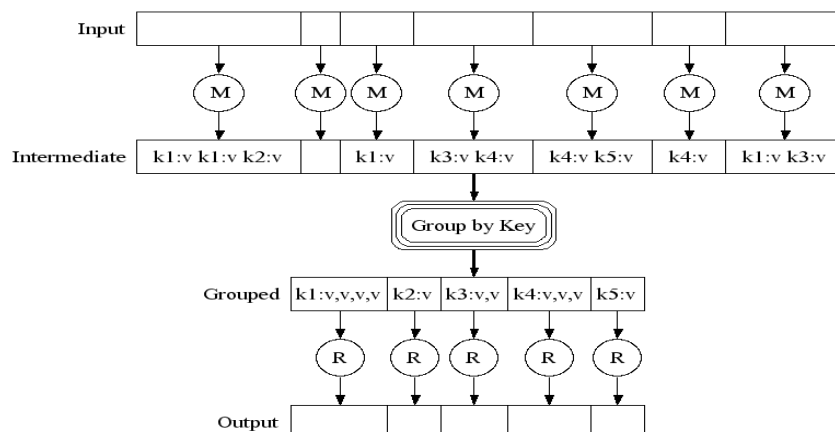
7

Example: Map Processing in Hadoop

- **Given a file**
 - A file may be divided into multiple parts (splits).
- **Each record (line) is processed by a Map function,**
 - written by the user,
 - takes an input key/value pair
 - produces a set of intermediate key/value pairs.
 - e.g. (doc—id, doc-content)
- **Draw an analogy to SQL *group-by* clause**

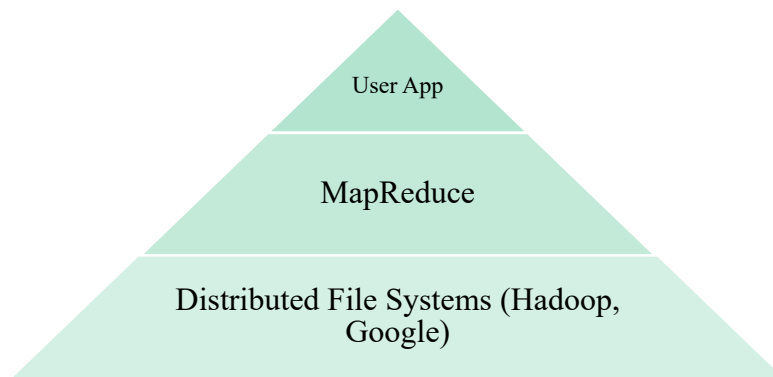
8

Put Map and Reduce Tasks Together



12

Systems Support for MapReduce



20

Distributed Filesystems

- **The interface is the same as a single-machine file system**
 - create(), open(), read(), write(), close()
- **Distribute file data to a number of machines (storage units).**
 - Support replication
- **Support concurrent data access**
 - Fetch content from remote servers. Local caching
- **Different implementations sit in different places on complexity/feature scale**
 - Google file system and Hadoop HDFS
 - » Highly scalable for large data-intensive applications.
 - » Provides redundant storage of massive amounts of data on cheap and unreliable computers

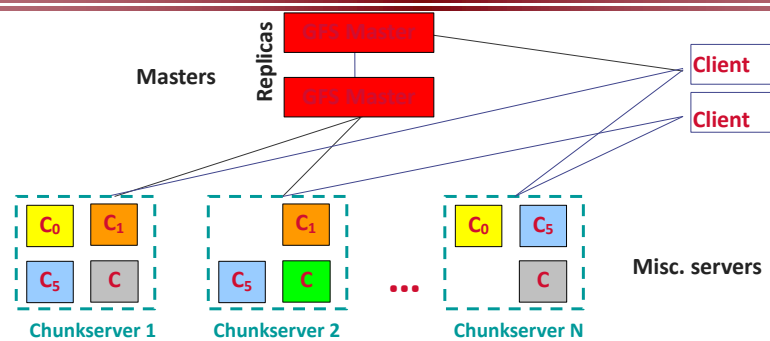
21

Assumptions of GFS/Hadoop DFS

- **High component failure rates**
 - Inexpensive commodity components fail all the time
- **“Modest” number of HUGE files**
 - Just a few million
 - Each is 100MB or larger; multi-GB files typical
- **Files are write-once, mostly appended to**
 - Perhaps concurrently
- **Large streaming reads**
- **High sustained throughput favored over low latency**

22

GFS Design




- Files are broken into chunks (typically 64 MB) and serve in chunk servers
- Master manages metadata, but clients may cache meta data obtained.
 - Data transfers happen directly between clients/chunk-servers
- Reliability through replication
 - Each chunk replicated across 3+ *chunk-servers*

23


Hadoop Distributed File System

- Files split into 128MB blocks
- Blocks replicated across several datanodes (often 3)
- Namenode stores metadata (file names, locations, etc)
- Optimized for large files, sequential reads
- Files are append-only



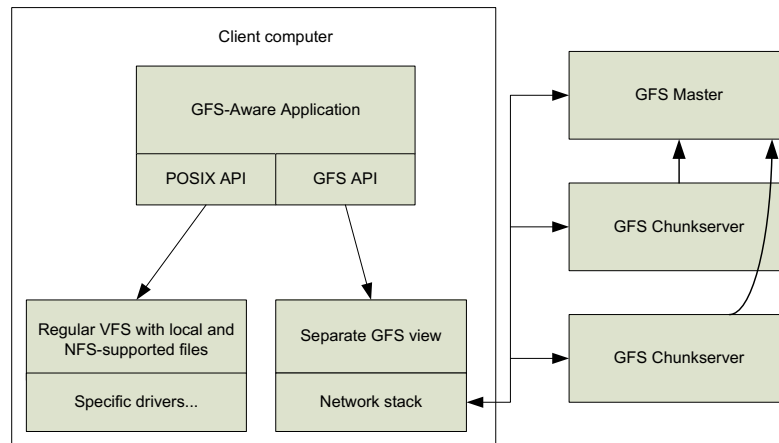
24

Hadoop DFS



25

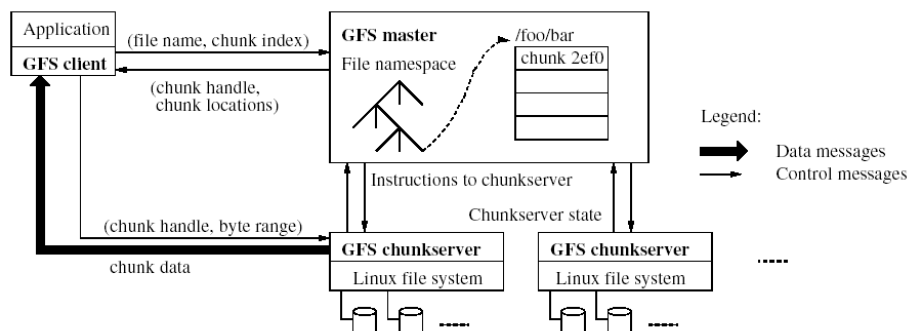
GFS Client Block Diagram



- Provide both POSIX standard file interface, and costumed API
- Can cache meta data for direct client-chunk server access

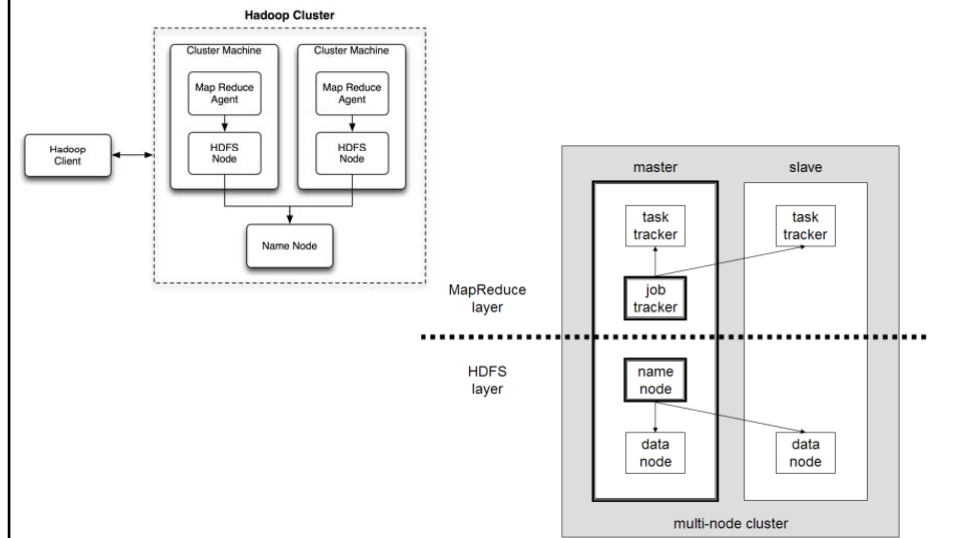
26

Read/write access flow in GFS



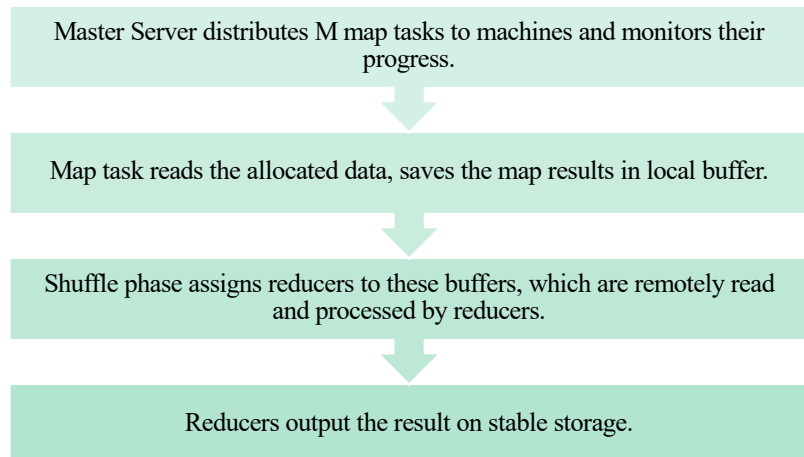
27

Hadoop DFS with MapReduce




28

MapReduce: Execution overview




29

Execute MapReduce on a cluster of machines with Hadoop DFS



30

MapReduce in Parallel: Example



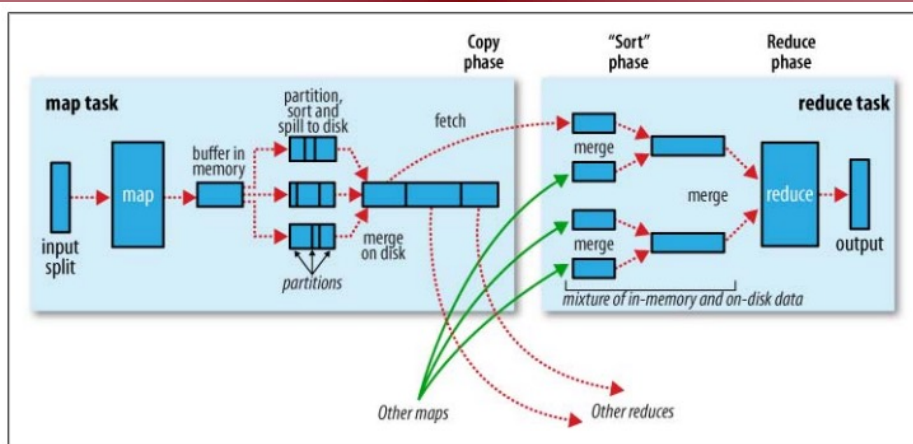
31

MapReduce: Execution Details

- **Input reader**
 - Divide input into splits, assign each split to a Map task
- **Map task**
 - Apply the Map function to each record in the split
 - Each Map function returns a list of (key, value) pairs
- **Shuffle/Partition and Sort**
 - Shuffle distributes sorting & aggregation to many reducers
 - All records for key k are directed to the same reduce processor
 - Sort groups the same keys together, and prepares for aggregation
- **Reduce task**
 - Apply the Reduce function to each key
 - The result of the Reduce function is a list of (key, value) pairs

32

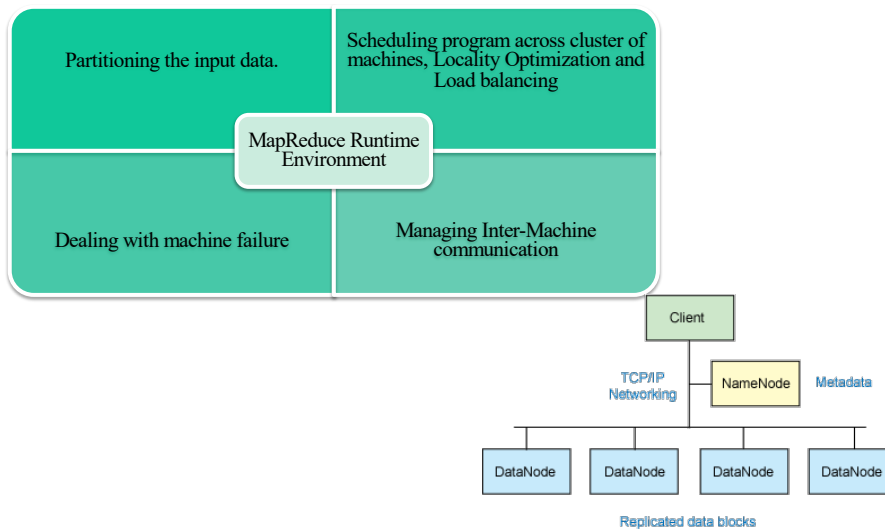
MapReduce with data shuffling & sorting



Tom White, *Hadoop: The Definitive Guide*

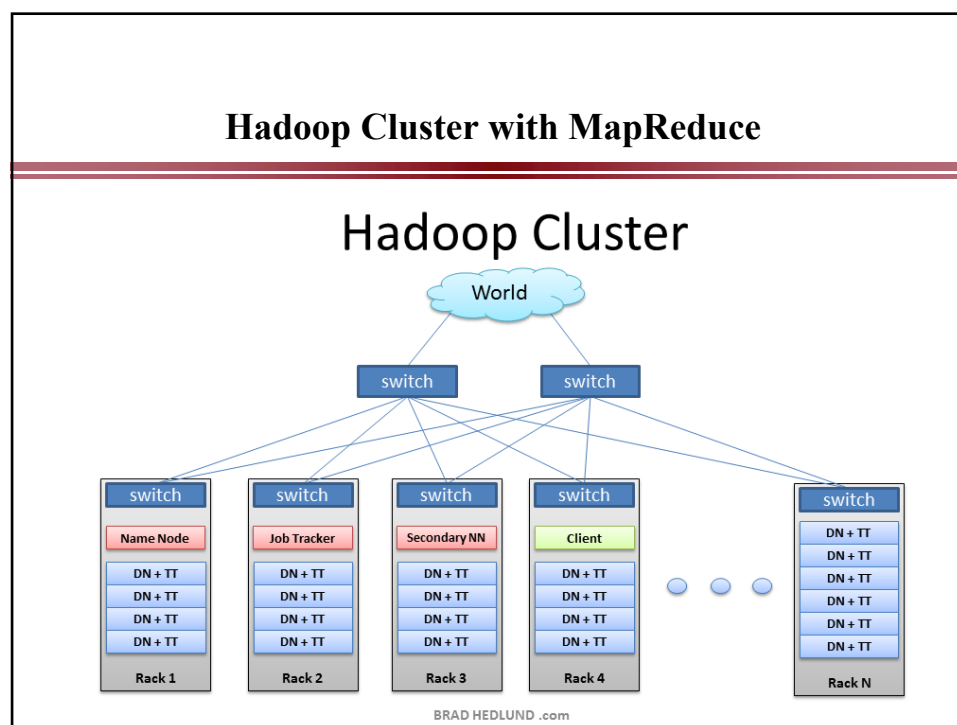
33

MapReduce: Runtime Environment & Hadoop



34

Hadoop Cluster with MapReduce



35

MapReduce: Fault Tolerance

- **Handled via re-execution of tasks.**
 - Task completion committed through master
- **Mappers save outputs to local disk before serving to reducers**
 - Allows recovery if a reducer crashes
 - Allows running more reducers than # of nodes
- **If a task crashes:**
 - Retry on another node
 - » OK for a map because it had no dependencies
 - » OK for reduce because map outputs are on disk
 - If the same task repeatedly fails, fail the job or ignore that input block
 - : For the fault tolerance to work, *user tasks must be deterministic and side-effect-free*
- 2. **If a node crashes:**
 - Relaunch its current tasks on other nodes
 - Relaunch any maps the node previously ran
 - » Necessary because their output files were lost along with the crashed node

36

MapReduce: Locality Optimization

- Leverage the distributed file system to schedule a map task on a machine that contains a replica of the corresponding input data.
- Thousands of machines read input at local disk speed
- Without this, rack switches limit read rate

37

MapReduce: Redundant Execution

- Slow workers are source of bottleneck, may delay completion time.
- Near end of phase, spawn backup tasks, one to finish first wins.
- Effectively utilizes computing power, reducing job completion time by a factor.

38

MapReduce: Skipping Bad Records

- Map/Reduce functions sometimes fail for particular inputs.
- Fixing the Bug might not be possible : Third Party Libraries.
- On Error
 - Worker sends signal to Master
 - If multiple error on same record, skip record

39

MapReduce: Miscellaneous Refinements

- Combiner function at a map task
- Sorting Guarantees within each reduce partition.
- Local execution for debugging/testing
- User-defined counters


40

Combining Phase

- Run on map machines after map phase
- “Mini-reduce,” only on local map output
- Used to save bandwidth before sending data to full reduce tasks
- Reduce tasks can be combiner if commutative & associative

41

Combiner, graphically




42

Examples of MapReduce Usage in Web Applications

- Distributed Grep.
- Count of URL Access Frequency.
- Clustering (K-means)
- Graph Algorithms.
- Indexing Systems

MapReduce Programs In Google Source Tree



43

Hadoop and Tools

- **Various Linux Hadoop clusters around**
 - Cluster +Hadoop
 - » <http://hadoop.apache.org>
 - Amazon EC2
- **Winows and other platforms**
 - The NetBeans plugin simulates Hadoop
 - The workflow view works on Windows
- **Hadoop-based tools**
 - For Developing in Java, NetBeans plugin
- **Pig Latin**, a SQL-like high level data processing script language
- **Hive**, Data warehouse, SQL
- **Mahout**, Machine Learning algorithms on Hadoop
- **HBase**, Distributed data store as a large table

44

More MapReduce Applications

- Map Only processing
- Filtering and accumulation
- Database join
- Reversing graph edges
- Producing inverted index for web search
- PageRank graph processing

45

MapReduce Use Case 1: Map Only

Data distributive tasks – Map Only

- E.g. classify individual documents
- Map does everything
 - Input: (docno, doc_content), ...
 - Output: (docno, [class, class, ...]), ...
- No reduce tasks

46

MapReduce Use Case 2: Filtering and Accumulation

Filtering & Accumulation – Map and Reduce

- E.g. Counting total enrollments of two given student classes
- Map selects records and outputs initial counts
 - In: (Jamie, 11741), (Tom, 11493), ...
 - Out: (11741, 1), (11493, 1), ...
- Shuffle/Partition by class_id
- Sort
 - In: (11741, 1), (11493, 1), (11741, 1), ...
 - Out: (11493, 1), ..., (11741, 1), (11741, 1), ...
- Reduce accumulates counts
 - In: (11493, [1, 1, ...]), (11741, [1, 1, ...])
 - Sum and Output: (11493, 16), (11741, 35)

47

MapReduce Use Case 3: Database Join

- A JOIN is a means for combining fields from two tables by using values common to each.
- Example :For each employee, find the department he works in

Employee Table		JOIN Pred: EMPLOYEE.DepID= DEPARTMENT.DepID	Department Table	
LastName	DepartmentID		DepartmentID	DepartmentName
Rafferty	31		31	Sales
Jones	33		33	Engineering
Steinberg	33		34	Clerical
Robinson	34		35	Marketing
Smith	34			

JOIN RESULT	
LastName	DepartmentName
Rafferty	Sales
Jones	Engineering
Steinberg	Engineering
...	...

48

MapReduce Use Case 3 – Database Join

Problem: Massive lookups

- Given two large lists: (URL, ID) and (URL, doc_content) pairs
- Produce (URL, ID, doc_content) or (ID, doc_content)

Solution:


- **Input stream:** both (URL, ID) and (URL, doc_content) lists
 - (http://del.icio.us/post, 0), (http://digg.com/submit, 1), ...
 - (http://del.icio.us/post, <html0>), (http://digg.com/submit, <html1>), ...
- **Map** simply passes input along,
- **Shuffle and Sort on URL** (group ID & doc_content for the same URL together)
 - Out: (http://del.icio.us/post, 0), (http://del.icio.us/post, <html0>), (http://digg.com/submit, <html1>), (http://digg.com/submit, 1), ...
- **Reduce** outputs result stream of (ID, doc_content) pairs
 - In: (http://del.icio.us/post, [0, html0]), (http://digg.com/submit, [html1, 1]), ...
 - Out: (0, <html0>), (1, <html1>), ...

49

MapReduce Use Case 4: Reverse graph edge directions & output in node order

- **Input example: adjacency list of graph (3 nodes and 4 edges)**

(3, [1, 2]) (1, [3])
 (1, [2, 3]) → (2, [1, 3])
 (3, [1])



- node_ids in the output **values** are also sorted.
 But Hadoop only sorts on keys!


- **MapReduce format**

50

MapReduce Use Case 4: Reverse graph edge directions & output in node order

- **Input example: adjacency list of graph (3 nodes and 4 edges)**

(3, [1, 2]) (1, [3])
 (1, [2, 3]) → (2, [1, 3])
 (3, [1])



- node_ids in the output **values** are also sorted.
 But Hadoop only sorts on keys!

- **MapReduce format**

- Input: (3, [1, 2]), (1, [2, 3]).
- Intermediate: (1, [3]), (2, [3]), (2, [1]), (3, [1]). (reverse edge direction)
- Out: (1,[3]) (2, [1, 3]) (3, [[1]).

51

MapReduce Use Case 5: Inverted Indexing Preliminaries

Construction of inverted lists for document search

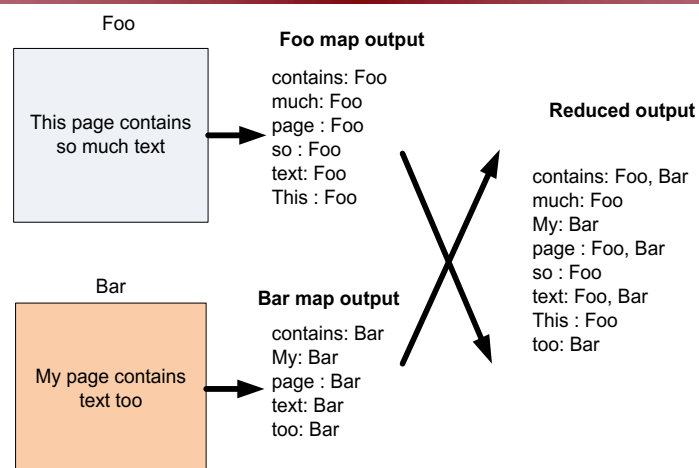
- Input: documents: (docid, [term, term..]), (docid, [term, ..]), ..
- Output: (term, [docid, docid, ...])
 - E.g., (apple, [1, 23, 49, 127, ...])

A document id is an internal document id, e.g., a unique integer

- Not an external document id such as a url


52

Inverted Index: Data flow



54

MapReduce Use Case 6: PageRank




60

PageRank

- Model page reputation on the web

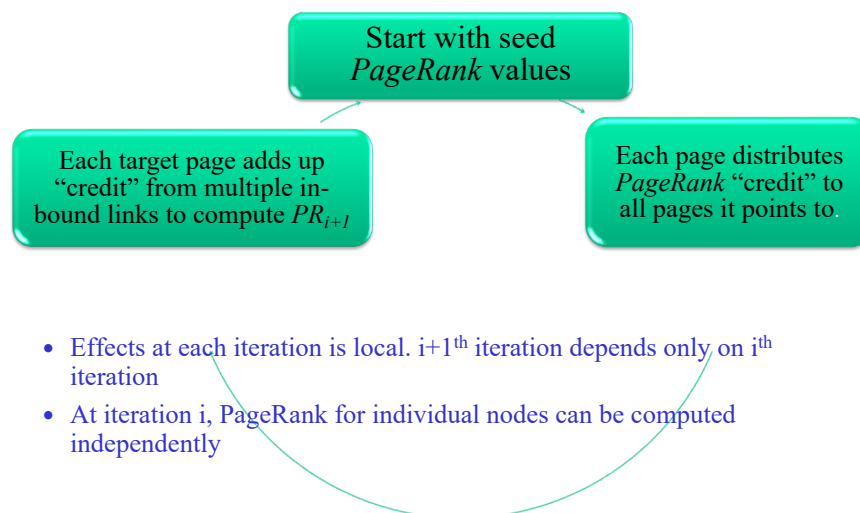
$$PR(x) = (1 - d) + d \sum_{i=1}^n \frac{PR(t_i)}{C(t_i)}$$

- $i=1, n$ lists all parents of page x .
- $PR(x)$ is the page rank of each page.
- $C(t)$ is the out-degree of t .
- d is a damping factor .



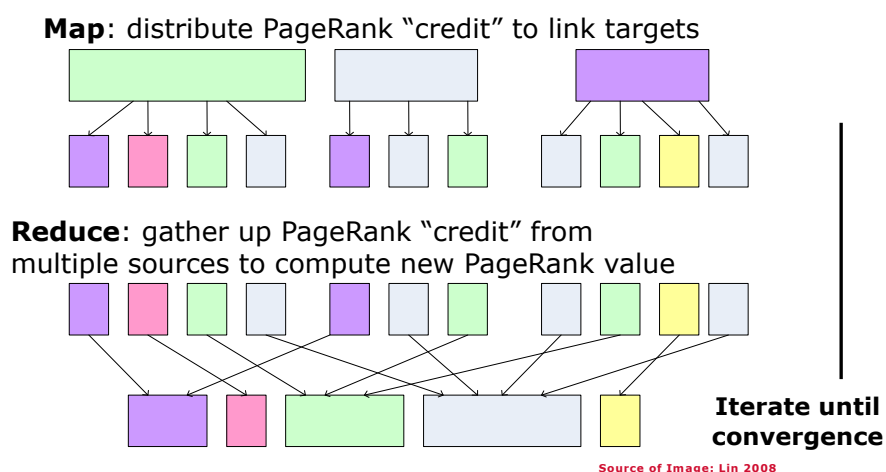
61

Computing PageRank Iteratively



62

PageRank using MapReduce



63

PageRank Calculation: Preliminaries

One PageRank iteration:

- Input:
 - $(id_1, [score_1^{(t)}, out_{11}, out_{12}, ..]), (id_2, [score_2^{(t)}, out_{21}, out_{22}, ..])$
 - ..
- Output:
 - $(id_1, [score_1^{(t+1)}, out_{11}, out_{12}, ..]), (id_2, [score_2^{(t+1)}, out_{21}, out_{22}, ..])$..

MapReduce elements

- Score distribution and accumulation
- Database join

64

PageRank: Score Distribution and Accumulation

- **Map**
 - In: $(id_1, [score_1^{(t)}, out_{11}, out_{12}, ..]), (id_2, [score_2^{(t)}, out_{21}, out_{22}, ..])$..
 - Out: $(out_{11}, score_1^{(t)}/n_1), (out_{12}, score_1^{(t)}/n_1) .., (out_{21}, score_2^{(t)}/n_2), ..$
- **Shuffle & Sort by node_id**
 - In: $(id_2, score_1), (id_1, score_2), (id_1, score_1), ..$
 - Out: $(id_1, score_1), (id_1, score_2), .., (id_2, score_1), ..$
- **Reduce**
 - In: $(id_1, [score_1, score_2, ..]), (id_2, [score_1, ..]), ..$
 - Out: $(id_1, score_1^{(t+1)}), (id_2, score_2^{(t+1)}), ..$

65

PageRank: Database Join to associate outlinks with score

- **Map**

- In & Out: $(id_1, score_1^{(t+1)})$, $(id_2, score_2^{(t+1)})$, .., $(id_1, [out_{11}, out_{12}, ..])$, $(id_2, [out_{21}, out_{22}, ..])$..

- **Shuffle & Sort by node_id**

- Out: $(id_1, score_1^{(t+1)})$, $(id_1, [out_{11}, out_{12}, ..])$, $(id_2, [out_{21}, out_{22}, ..])$, $(id_2, score_2^{(t+1)})$, ..

- **Reduce**

- In: $(id_1, [score_1^{(t+1)}, out_{11}, out_{12}, ..])$, $(id_2, [out_{21}, out_{22}, .., score_2^{(t+1)}])$, ..
- Out: $(id_1, [score_1^{(t+1)}, out_{11}, out_{12}, ..])$, $(id_2, [score_2^{(t+1)}, out_{21}, out_{22}, ..])$..

66

Conclusions

- **MapReduce advantages**

- **Application cases**

- Map only: for totally distributive computation
- Map+Reduce: for filtering & aggregation
- Database join: for massive dictionary lookups
- Secondary sort: for sorting on values
- Inverted indexing: combiner, complex keys
- PageRank: side effect files

67

For More Information

- J. Dean and S. Ghemawat. “MapReduce: Simplified Data Processing on Large Clusters.” *Proceedings of the 6th Symposium on Operating System Design and Implementation (OSDI 2004)*, pages 137-150. 2004.
- S. Ghemawat, H. Gobioff, and S.-T. Leung. “The Google File System.” *OSDI 2000?*
- http://hadoop.apache.org/common/docs/current/mapred_tutorial.html. “Map/Reduce Tutorial”. Fetched January 21, 2010.
- Tom White. *Hadoop: The Definitive Guide*. O'Reilly Media. June 5, 2009
- <http://developer.yahoo.com/hadoop/tutorial/module4.html>
- J. Lin and C. Dyer. *Data-Intensive Text Processing with MapReduce*. Book Draft. February 7, 2010